

Chapter 4

Network Layer


A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- ❖ If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2012
J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

6th edition

Jim Kurose, Keith Ross
Addison-Wesley

March 2012

Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

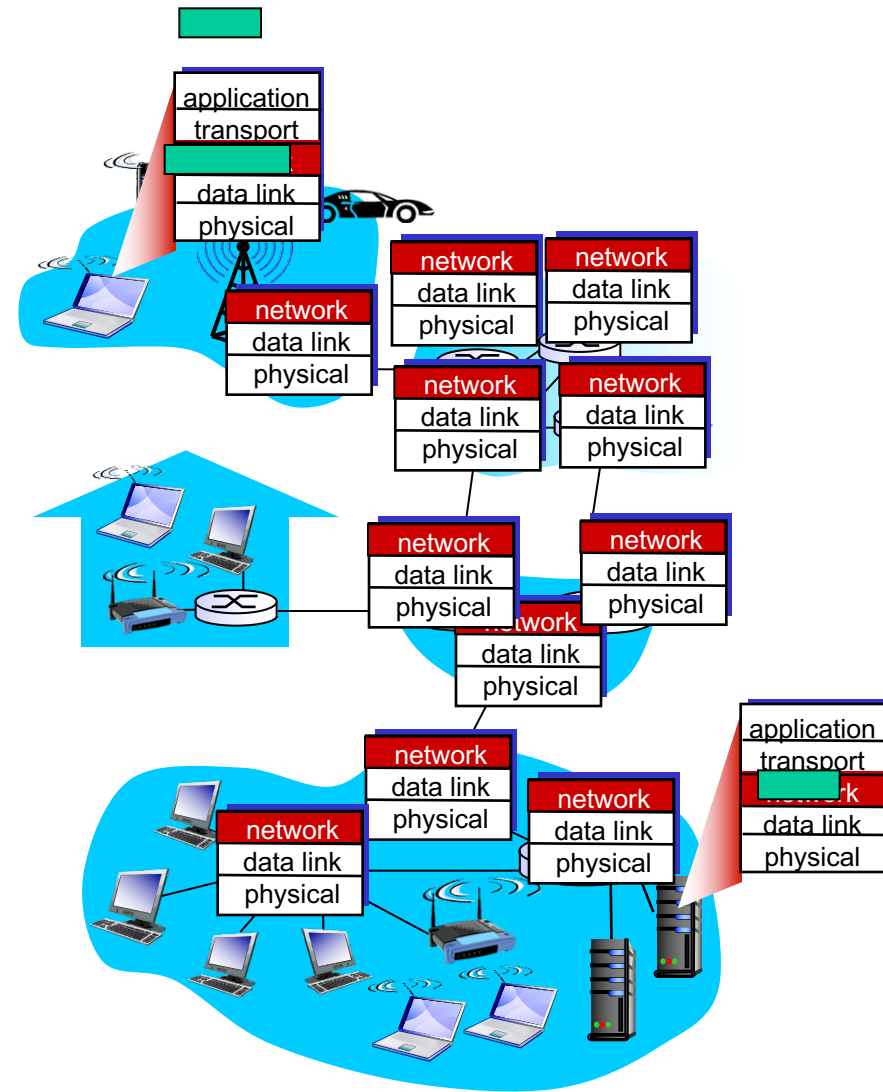
- RIP
- OSPF
- BGP

Link layer issues

- ❖ inefficient link-redundancy management
 - Spanning-tree Protocol
- ❖ filtering-database saturation on switches in large-size networks
 - route aggregation not possible on switches
- ❖ layer-2 broadcast traffic propagation
 - broadcast traffic is useful for several application but *must* be limited

Network layer

- ❖ transport segment from sending to receiving host
- ❖ on sending side encapsulates segments into datagrams
- ❖ on receiving side, delivers segments to transport layer
- ❖ network layer protocols in *every* host, router
- ❖ router examines header fields in all IP datagrams passing through it



Two key network-layer functions

- ❖ *forwarding*: move packets from router's input to appropriate router output

- ❖ *routing*: determine route taken by packets from source to dest.

 - *routing algorithms*

analogy:

- ❖ *routing*: process of planning trip from source to dest

- ❖ *forwarding*: process of getting through single interchange

Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

- RIP
- OSPF
- BGP

Connection, connection-less service

- ❖ *datagram* network provides network-layer *connectionless* service
- ❖ *virtual-circuit (VC)* network provides network-layer *connection* service

Virtual circuits

“source-to-dest path behaves much like telephone circuit”

- performance-wise
- network actions along source-to-dest path

- ❖ call setup, teardown for each call *before* data can flow
- ❖ each packet carries VC identifier (not destination host address)
- ❖ every router on source-dest path maintains “state” for each passing connection
- ❖ link, router resources (bandwidth, buffers) may be *allocated* to VC (dedicated resources = predictable service)


VC implementation

a VC consists of:

1. *path* from source to destination
 2. *VC numbers*, one number for each link along path
 3. *entries in forwarding tables* in routers along path
- ❖ packet belonging to VC carries VC number (rather than dest address)
 - ❖ VC number can be changed on each link.
 - new VC number comes from forwarding table

VC forwarding table

*forwarding table in
northwest router:*




| Incoming interface | Incoming VC # | Outgoing interface | Outgoing VC # |
|--------------------|---------------|--------------------|---------------|
| 1 | 12 | 3 | 22 |
| 2 | 63 | 1 | 18 |
| 3 | 7 | 2 | 17 |
| 1 | 97 | 3 | 87 |
| ... | ... | ... | ... |

VC routers maintain connection state information!


Virtual circuits: signaling protocols

- ❖ used to setup, maintain teardown VC
- ❖ used in ATM, frame-relay, X.25
- ❖ not used in today's Internet

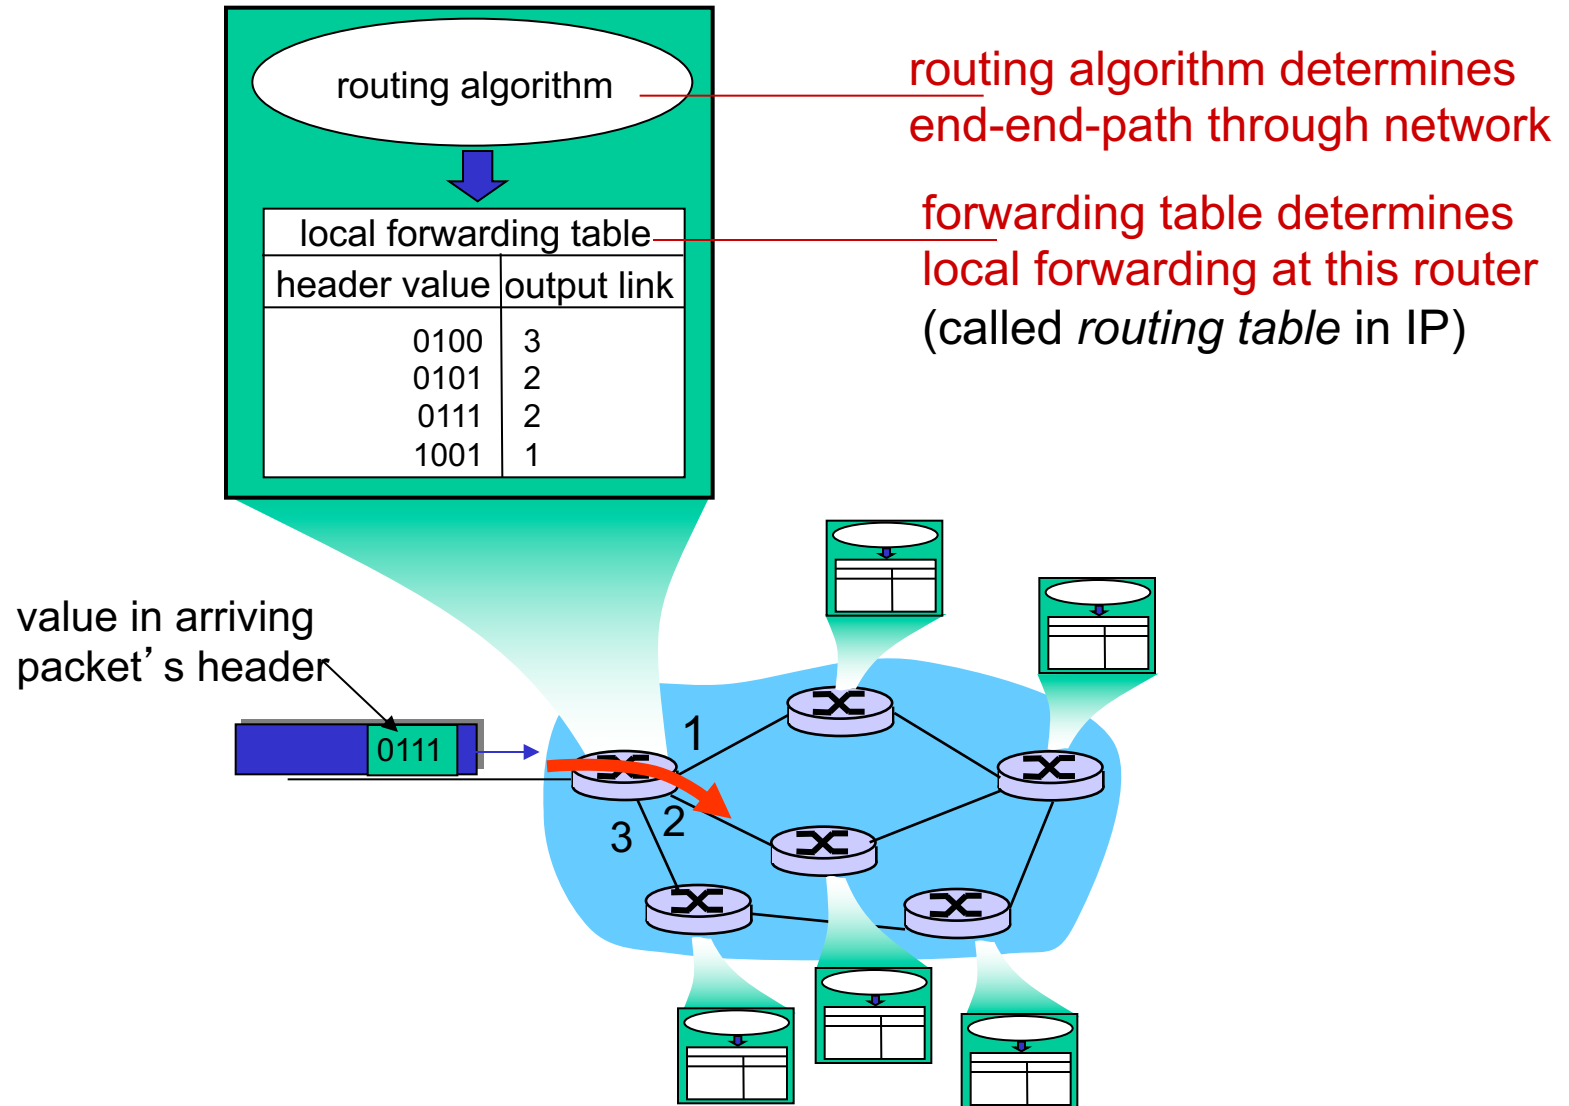


Datagram networks

- ❖ no call setup at network layer
- ❖ routers: no state about end-to-end connections
 - no network-level concept of “connection”
- ❖ packets forwarded using destination host address



Datagram forwarding table



Datagram or VC network: why?

Internet (datagram)

- ❖ data exchange among computers
 - “elastic” service, no strict timing req.
- ❖ many link types
 - different characteristics
 - uniform service difficult
- ❖ “smart” end systems (computers)
 - can adapt, perform control, error recovery
 - ***simple inside network, complexity at “edge”***

ATM (VC)

- ❖ evolved from telephony
- ❖ human conversation:
 - strict timing, reliability requirements
 - need for guaranteed service
- ❖ “dumb” end systems
 - telephones
 - ***complexity inside network***

Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

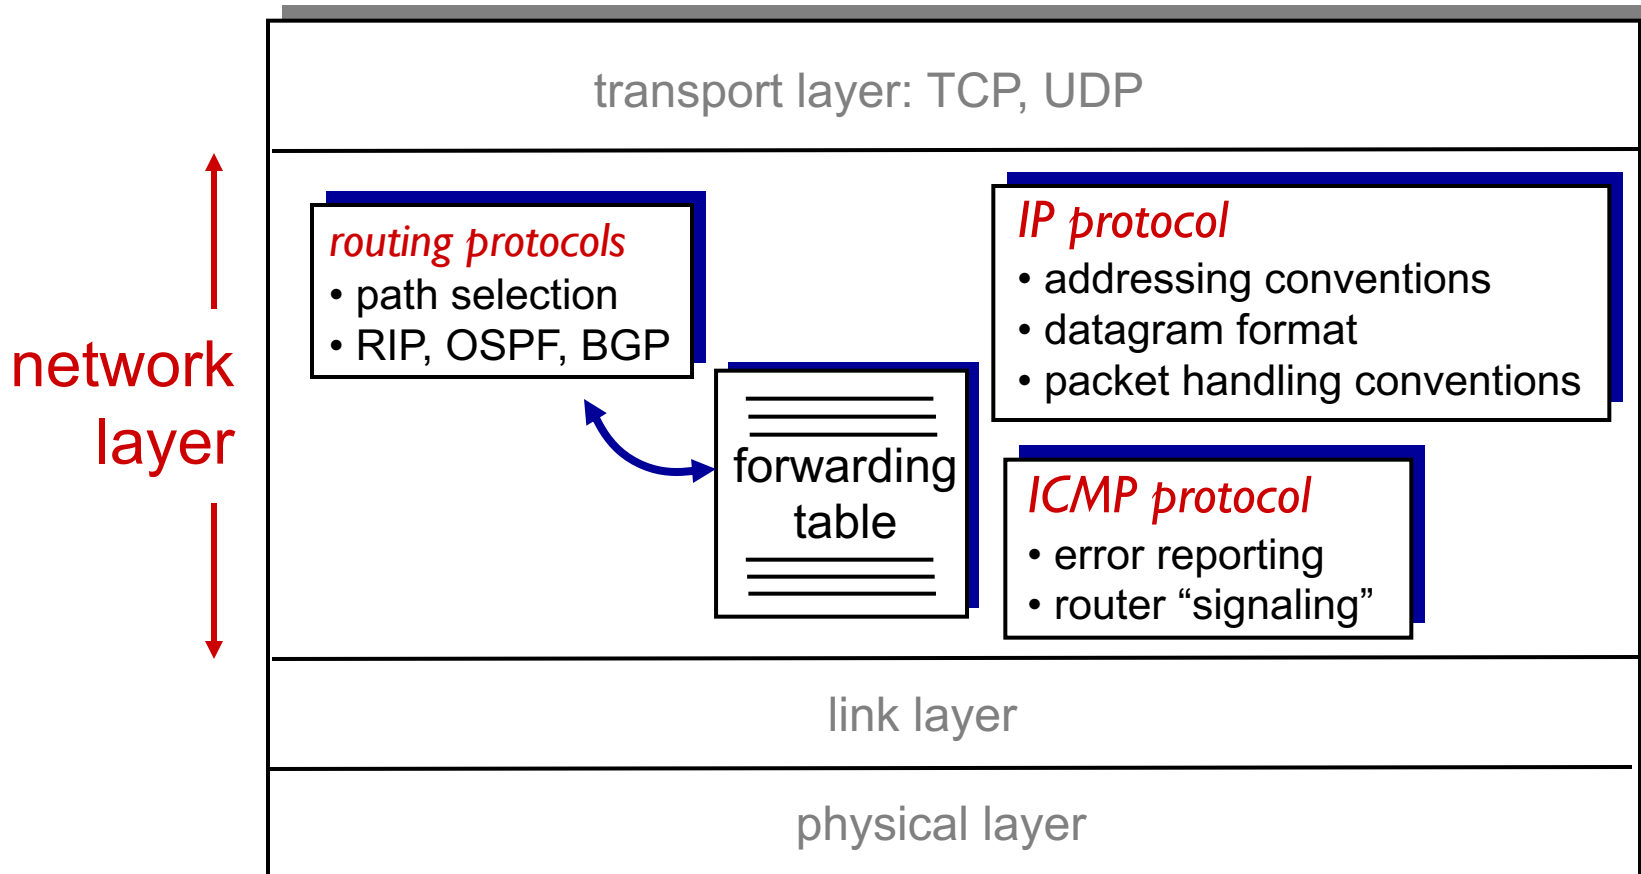
- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

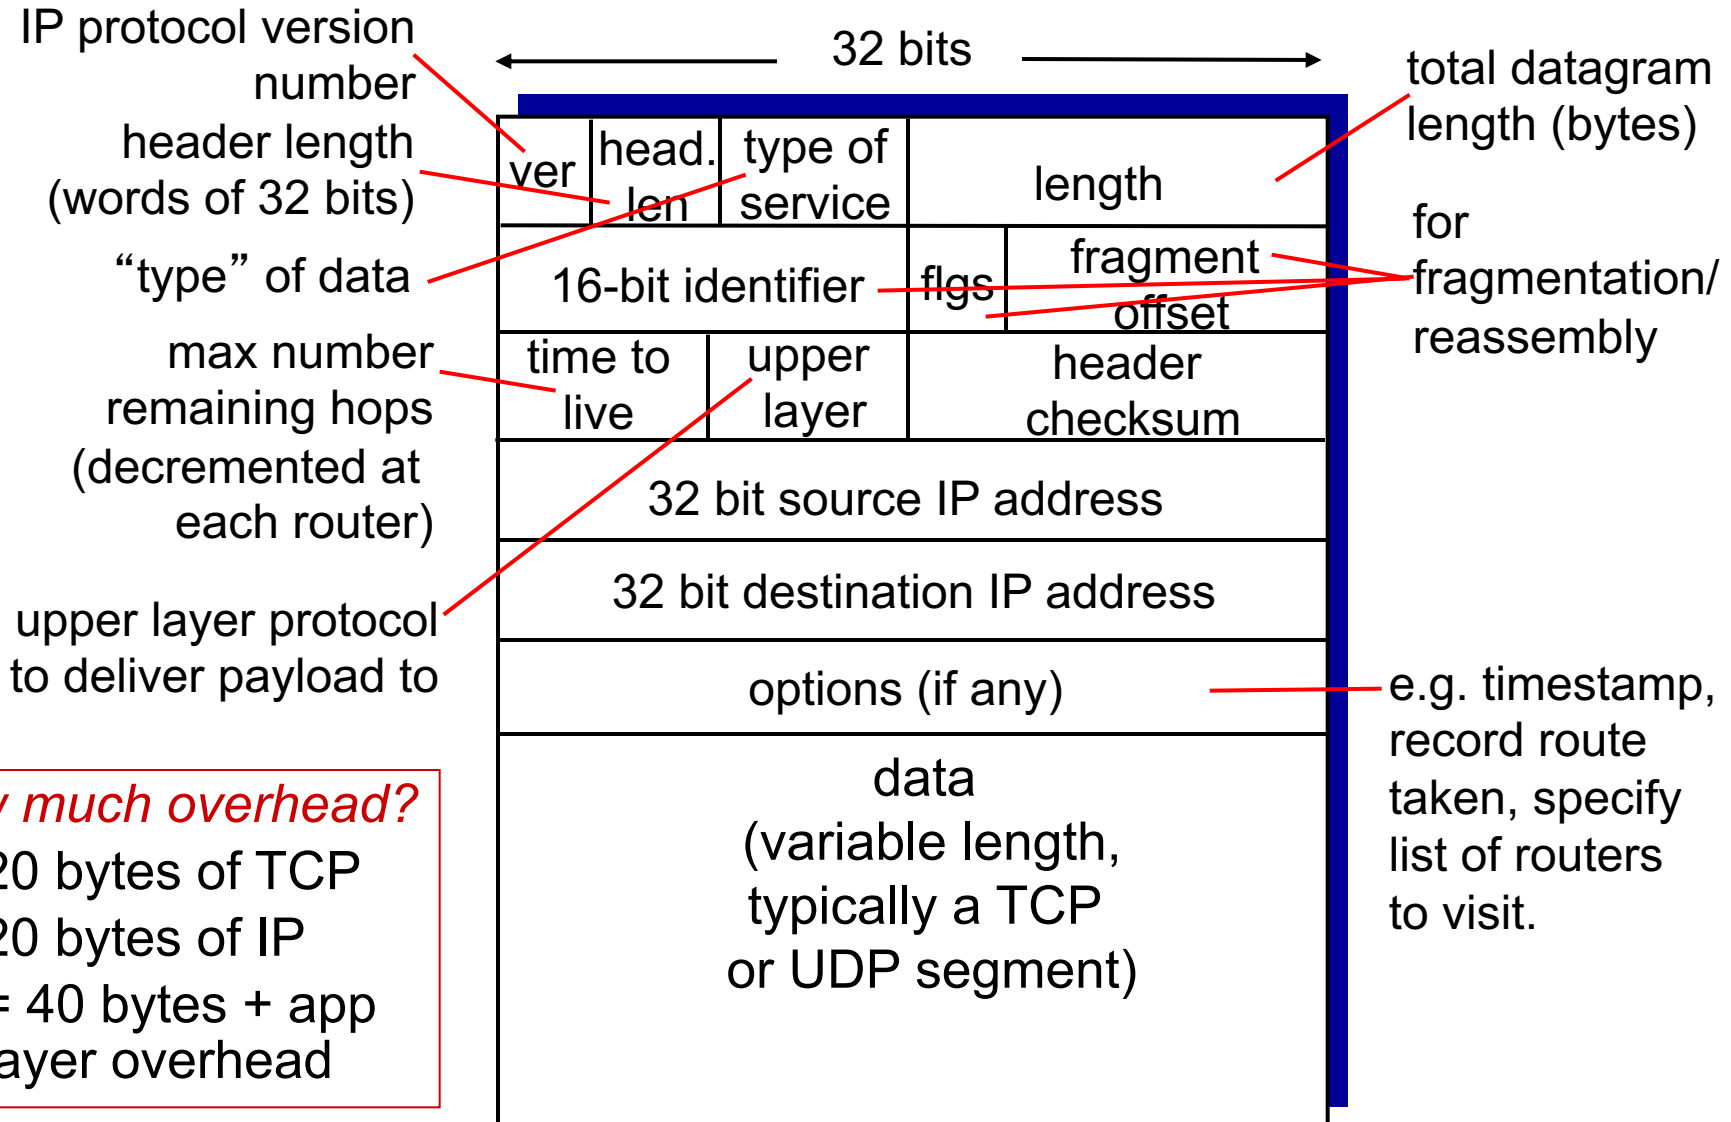
- RIP
- OSPF
- BGP

The Internet network layer

host, router network layer functions:



IP datagram format

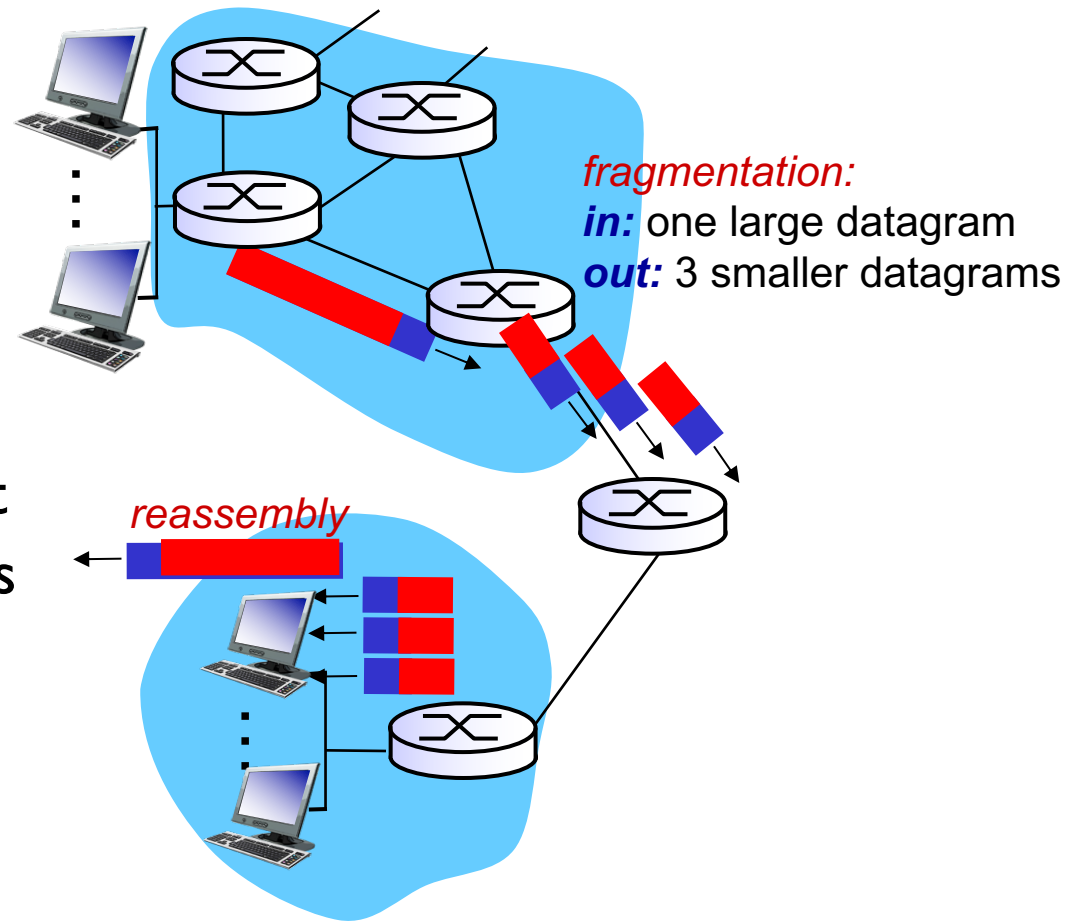


Header length vs. total length

- ❖ *header length* (4 bits) provides the number of 32 bit words forming the IP header
 - required due to the possible presence of IP options
 - e.g., 20 bytes + two 32 bits options \rightarrow HL = 7
- ❖ *total length* (16 bits) provides the total size of the IP datagram, including the header
 - maximum IP datagram size is $2^{16} - 1 = 65535$ bytes
 - required for properly handling fragmentation
 - also useful when padding at layer 2 is used
 - e.g., 46 bytes of Ethernet payload, TL = 40 \rightarrow padding = 6 bytes

IP fragmentation, reassembly

- ❖ network links have MTU (max.transfer size) - largest possible link-level frame
 - different link types, different MTUs
- ❖ large IP datagram divided (“fragmented”) within net
 - one datagram becomes several datagrams
 - “reassembled” only at final destination
 - IP header bits used to identify, order related fragments



IP fragmentation, reassembly

example:

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

| | length | ID | fragflag | offset | |
|--|--------|----|----------|--------|--|
| | =4000 | =x | =000 | =0 | |

*one large datagram becomes
several smaller datagrams*

1480 bytes in
data field

offset =
1480/8

| | length | ID | fragflag | offset | |
|--|--------|----|----------|--------|--|
| | =1500 | =x | =001 | =0 | |

| | length | ID | fragflag | offset | |
|--|--------|----|----------|--------|--|
| | =1500 | =x | =001 | =185 | |

| | length | ID | fragflag | offset | |
|--|--------|----|----------|--------|--|
| | =1040 | =x | =000 | =370 | |

flags (3 bits)

- ❖ 1st reserved (set to 0)
- ❖ 2nd set to 1 if don't fragment
- ❖ 3rd set to 1 if more fragments

offset: position (in multiple of 8 bytes) of
the fragment in the original datagram

Header checksum

- ❖ calculated considering all the fields of the IP header
- ❖ calculated also at the destination and compared with the value carried in the datagram
 - if same value, datagram delivered to upper layers
 - if not the same, datagram considered corrupted
- ❖ re-evaluated at each router, as also TTL changes

Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms


- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

- RIP
- OSPF
- BGP

IP addressing: introduction

- ❖ **IP address:** 32-bit identifier for host, router interface
- ❖ **interface:** connection between host/router and physical link
 - router's typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- ❖ **IP addresses associated with each interface**



$$223.1.1.1 = \underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$$

IP addressing: introduction


Q: *how are interfaces actually connected?*

A: *by means of a proper link layer technology.*

A: wired Ethernet interfaces connected by Ethernet switches

IP does not care about how one interface is connected to another (with no intervening router)

A: wireless WiFi interfaces connected by WiFi base station




Logical IP Subnets (LIS)

❖ IP address:

- subnet part - high order bits
- host part - low order bits

❖ *what 's a subnet ?*


- set of device interfaces with same subnet part of IP address
- interfaces *should* physically reach each other. *Why?*



network consisting of 3 subnets


Physical networks

- ❖ set of devices that can physically reach each other *by means of link layer mechanisms*
- ❖ to determine the physical networks, detach each interface from its router, creating islands of isolated networks



Subnets and physical networks


- ❖ IP assumes a *biunivocal correspondence* between physical networks and subnets
 - link layer mechanisms within the same subnet (*direct communication*)
 - IP routing between different subnets (*indirect communication*)



Subnets and physical networks


- ❖ more subnets over a single physical network are possible

- *one-arm router*



- ❖ one subnet over more physical networks is possible

- *proxy ARP*



IP addressing: special addresses

| | |
|--------|--------|
| Subnet | All 0s |
|--------|--------|

the (sub)network (network ID)

| |
|--------|
| All 1s |
|--------|

limited broadcast (local net)

| | |
|--------|--------|
| Subnet | All 1s |
|--------|--------|

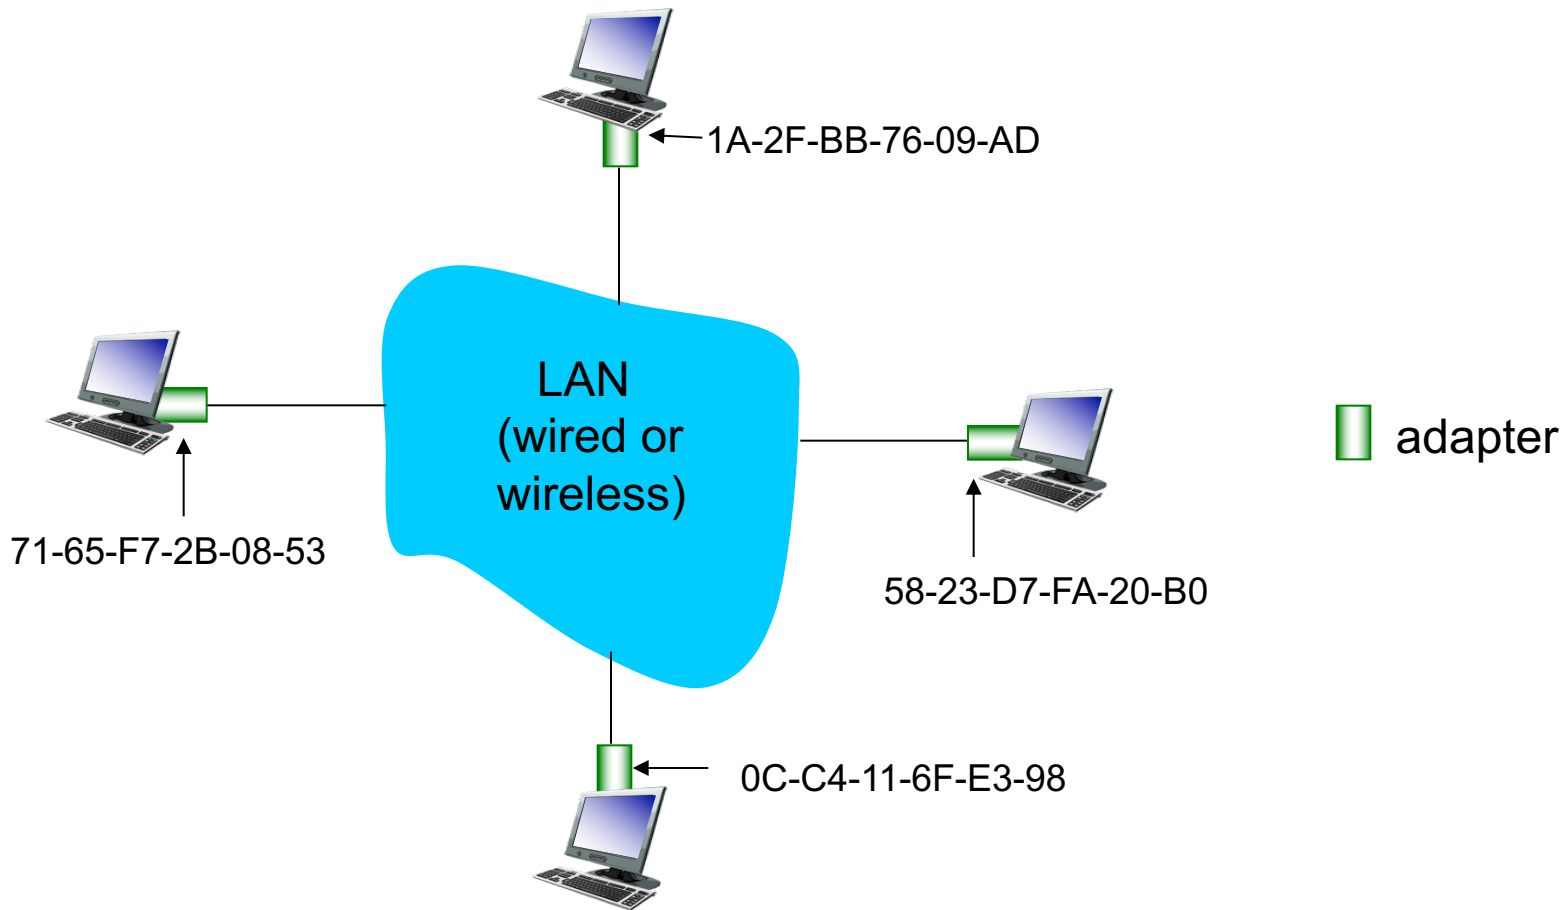
directed broadcast for net

| | |
|-----|--------------------|
| 127 | Anything (often 1) |
|-----|--------------------|

loopback

LAN addresses

each adapter on LAN has unique **LAN** address




LAN addresses (more)

- ❖ MAC address allocation administered by IEEE
- ❖ manufacturer buys portion of MAC address space (to assure uniqueness)
- ❖ analogy:
 - MAC address: like Social Security Number
 - IP address: like postal address
- ❖ MAC flat address → portability
 - can move LAN card from one LAN to another
- ❖ IP hierarchical address *not* portable
 - address *depends on IP subnet* to which node is attached

Subnets and IP routing


- ❖ IP addresses grouped according to their location (subnet)
- ❖ an entry on a routing table can refer to *an entire subnet* rather than to a single address!
- ❖ how can we do even better?
 - hierarchical addressing (see *later...*)



| Router A routing table | |
|------------------------|-------------|
| destination | output link |
| 223.1.1.0 | 1 |
| 223.1.2.0 | 2 |
| 223.1.3.0 | 2 |

ARP: address resolution protocol

Question: how to determine interface's MAC address, knowing its IP address?



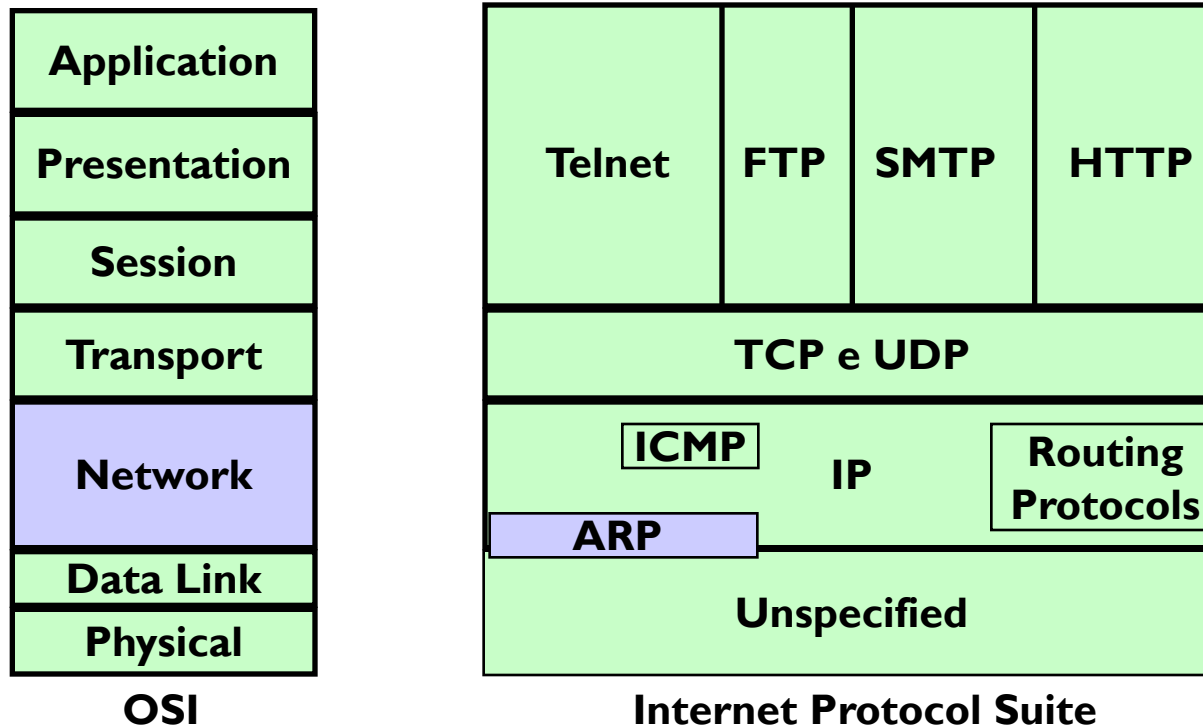
ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)

ARP protocol: same LAN

- ❖ A wants to send datagram to B
 - B's MAC address not in A's ARP table.
- ❖ A **broadcasts** ARP query packet (**ARP request**), containing B's IP address
 - dest MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive ARP query
- ❖ B receives ARP packet, replies to A (**ARP reply**) with its (B's) MAC address
 - frame sent to A's MAC address (unicast)
- ❖ A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)
 - soft state: information that times out (goes away) unless refreshed
- ❖ ARP is “plug-and-play”:
 - nodes create their ARP tables *without intervention from net administrator*

ARP protocol (more)



- ❖ ARP packets are not IP datagrams!
 - They do not have IP source/destination, TTL, etc.
- ❖ ARP packets are payloads of link layer frames

ARP request and reply

most significant fields of the link layer header

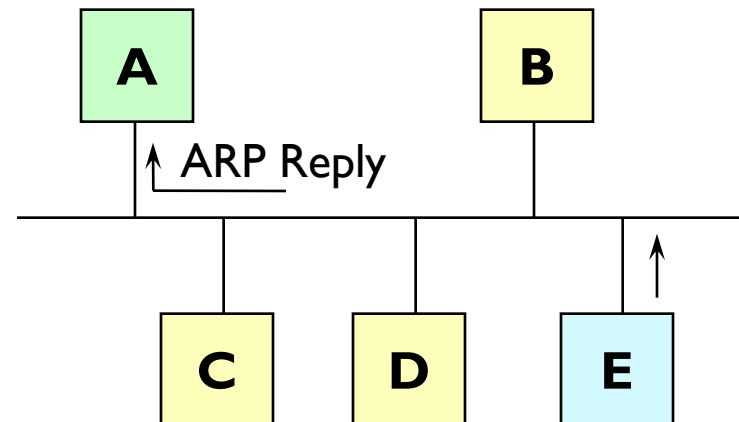
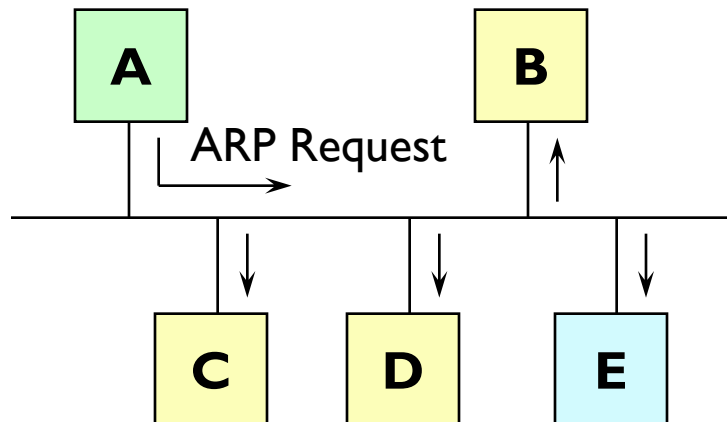
most significant fields of the ARP packet

| | | | | | | |
|----------------------|--------------|----------------|--------------|--------------|-----------|--------------|
| MAC broadcast | MAC A | ARP Req | MAC A | IP A* | ?? | IP E* |
|----------------------|--------------|----------------|--------------|--------------|-----------|--------------|

ARP Request

| | | | | | | |
|--------------|--------------|------------------|--------------|--------------|--------------|--------------|
| MAC A | MAC E | ARP Reply | MAC E | IP E* | MAC A | IP A* |
|--------------|--------------|------------------|--------------|--------------|--------------|--------------|

ARP Reply




*These fields are not IP source and destination fields of an IP header!

Addressing: routing to another LAN


walkthrough: **send datagram from A to B via R**

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)



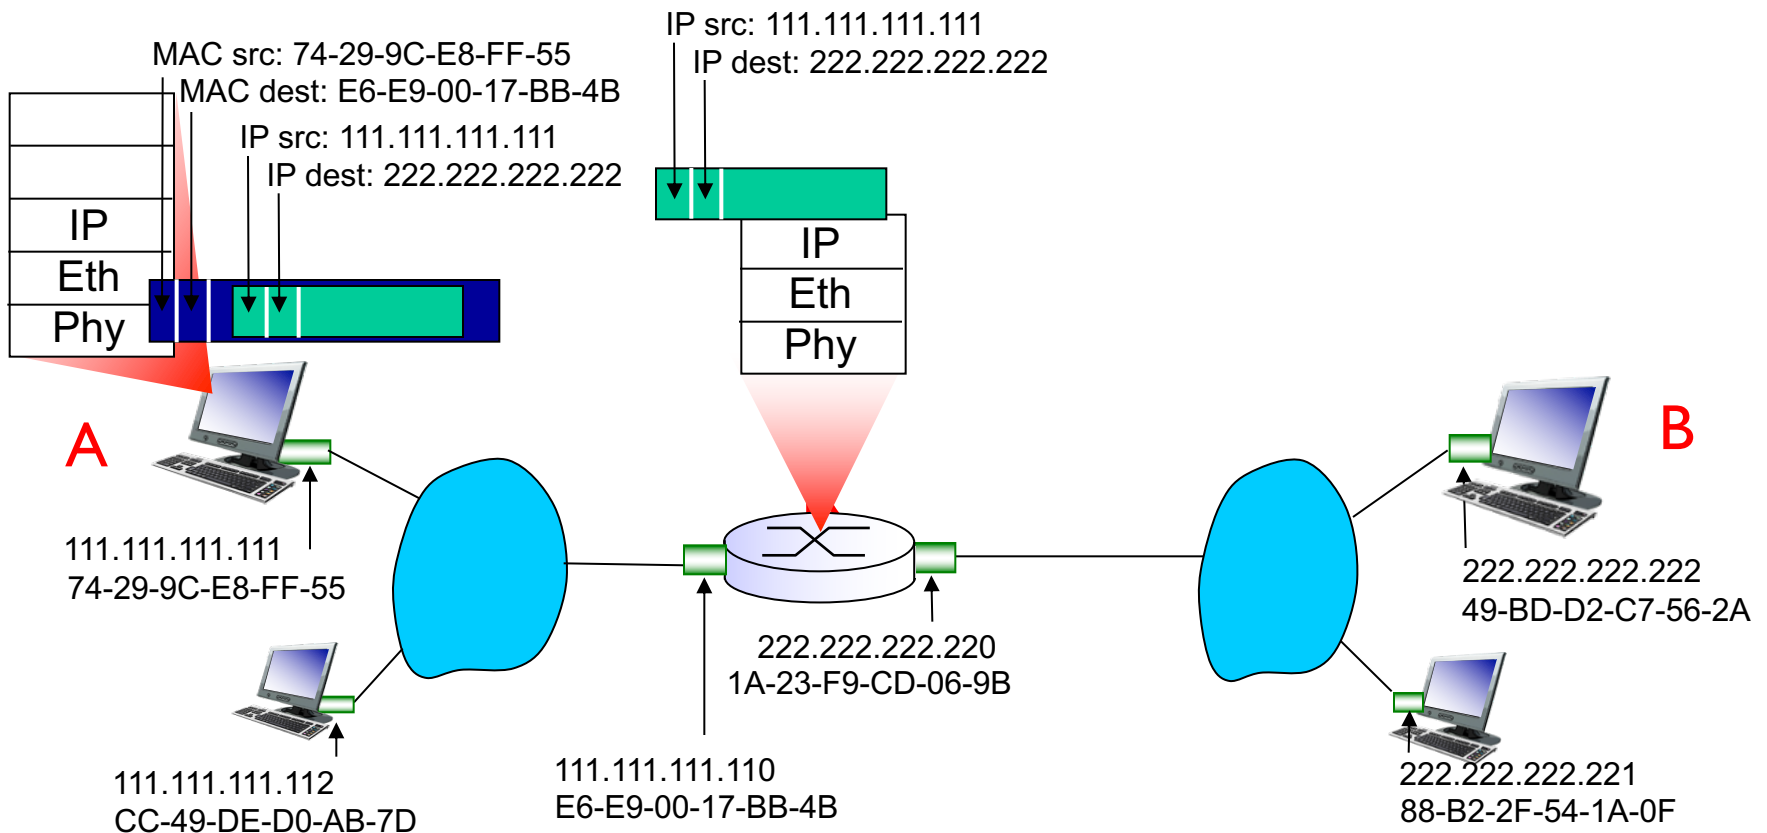
Addressing: routing to another LAN

- ❖ A creates IP datagram with IP source A, destination B
- ❖ A creates link-layer frame with R's MAC address as dest, frame contains A-to-B IP datagram




Addressing: routing to another LAN

- ❖ frame sent from A to R
- ❖ frame received at R, datagram removed, passed up to IP




Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram




IP addressing: the entire story...

- ❖ *classful addressing*: static division of network part and host part of an IP address
 - no concept of subnet
 - only three possible sizes for IP networks
 - poor flexibility
- ❖ *subnetting*: starting from a classful addressing, define smaller IP networks called “subnets”
 - introduction of the concept of subnet and subnet mask
 - Variable Length Subnet Masking (VLSM)
- ❖ *classless addressing*: concept of IP class completely removed
 - we should refer again to IP networks (and netmasks) rather than to IP subnets (and subnet masks), but nobody cares about it!
 - Classless Inter Domain Routing (CIDR)

IP addressing: Classes



Class A – 128 networks – 1st byte: 0-127



Class B – 16K networks – 1st byte: 128-191




Class C – 2M networks – 1st byte: 192-223

IP addressing: Classes



Class D – 1^o byte: 224-239




Class E – 1^o byte: 240-255

IP addressing: subnetting

VLSM: Variable Length Subnet Masking

- starting from an IP network of a given class, define smaller subnets of arbitrary size
- address format: **network ID** + **subnet mask**
 - **subnet mask**: all '1s' in the subnet part, all '0s' in the host part




136.23.16.0 255.255.254.0 ← subnet mask notation

IP addressing: CIDR

CIDR: Classless InterDomain Routing

- network portion of address of arbitrary length
- address format: **network ID** + **prefix length** or **netmask**
 - **prefix length**: /x, where x is # bits in network portion of address
 - **netmask**: all '1s' in the network part, all '0s' in the host part



200.23.16.0/23 ← **prefix length notation**

200.23.16.0 255.255.254.0 ← **netmask notation**

IP addressing: CIDR

valid netmasks: possible values in the 4 bytes composing the address

| | | |
|-----|-----------|-------|
| 0 | 0000 0000 | (256) |
| 128 | 1000 0000 | (128) |
| 192 | 1100 0000 | (64) |
| 224 | 1110 0000 | (32) |
| 240 | 1111 0000 | (16) |
| 248 | 1111 1000 | (8) |
| 252 | 1111 1100 | (4) |
| 254 | 1111 1110 | (2)* |
| 255 | 1111 1111 | (1) |

*not valid in the 4^o byte


IP addressing: CIDR

some examples

- 130.192.0.0/16 – 130.192.0.0 255.255.0.0
- 130.192.0.0/24 – 130.192.0.0 255.255.255.0
- 130.192.0.0/25 – 130.192.0.0 255.255.255.128
- 130.192.2.0/23 – 130.192.2.0 255.255.254.0
- 130.192.1.4/30 – 130.192.1.4 255.255.255.252
- ~~130.192.1.0/31 – 130.192.1.0 255.255.255.254~~

Each IP network *must* contain at least the network ID and the broadcast address!

IP addressing: a real example



IP addressing: device config

each device must be provided with

- ❖ *an IP address*
- ❖ *a netmask*
 - to *evaluate* its own network ID
 - to *infer* the network ID of the destination
- ❖ *a default gateway*, a.k.a. first-hop router
 - to perform indirect communications

Hierarchical addressing plans

Q: how does *organizations* get network part of IP addr?

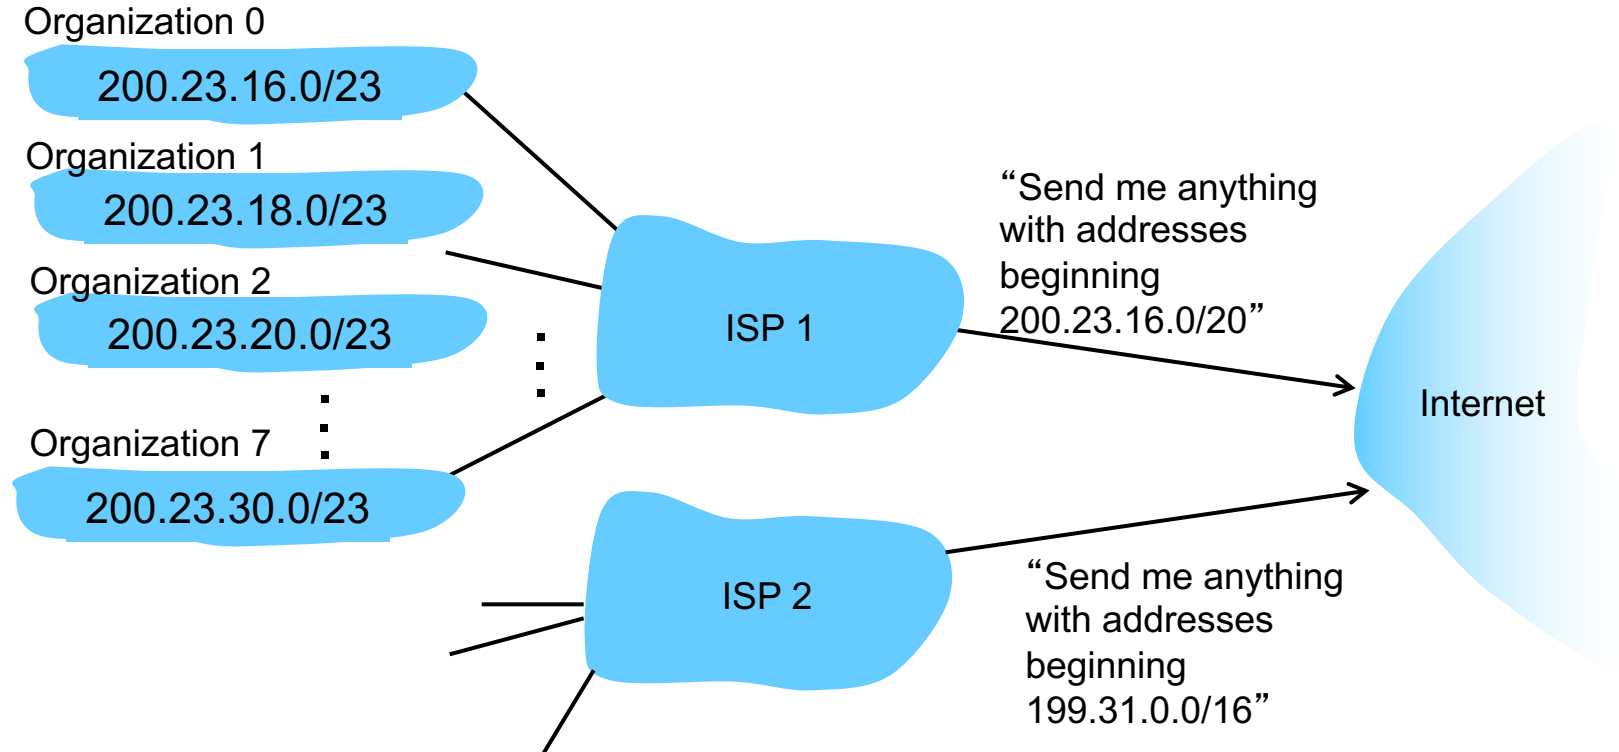
A: gets allocated portion of its provider ISP' s address space

| | | | | | |
|----------------|-----------------|-----------------|-----------------|----------|----------------|
| ISP's block | <u>11001000</u> | <u>00010111</u> | <u>00010000</u> | 00000000 | 200.23.16.0/20 |
| Organization 0 | <u>11001000</u> | <u>00010111</u> | <u>00010000</u> | 00000000 | 200.23.16.0/23 |
| Organization 1 | <u>11001000</u> | <u>00010111</u> | <u>00010010</u> | 00000000 | 200.23.18.0/23 |
| Organization 2 | <u>11001000</u> | <u>00010111</u> | <u>00010100</u> | 00000000 | 200.23.20.0/23 |
| ... | | | | | |
| Organization 7 | <u>11001000</u> | <u>00010111</u> | <u>00011110</u> | 00000000 | 200.23.30.0/23 |

*This **was not** possible with classful addressing!*

Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:



How do IP routing really work?

Q: How do a host learn its own IP network(s)?

❖ *bit-wise AND* between its IP address(es) and its netmask(s)

interface IP address: 192.168.10.69

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

bit-wise AND result: 192.168.10.64

netmask: 255.255.255.192

How do IP routing really work?

Q: How do a host infer the IP network of the destination it wants to contact?

❖ *bit-wise AND* between the destination IP address and *its own* netmask

destination IP address: 192.168.10.101

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

bit-wise AND result: 192.168.10.64

netmask: 255.255.255.192

How do IP routing really work?

Q: How do a host infer the IP network of the destination it wants to contact?

❖ *bit-wise AND* between the destination IP address and *its own* netmask

destination IP address: 192.168.10.132

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

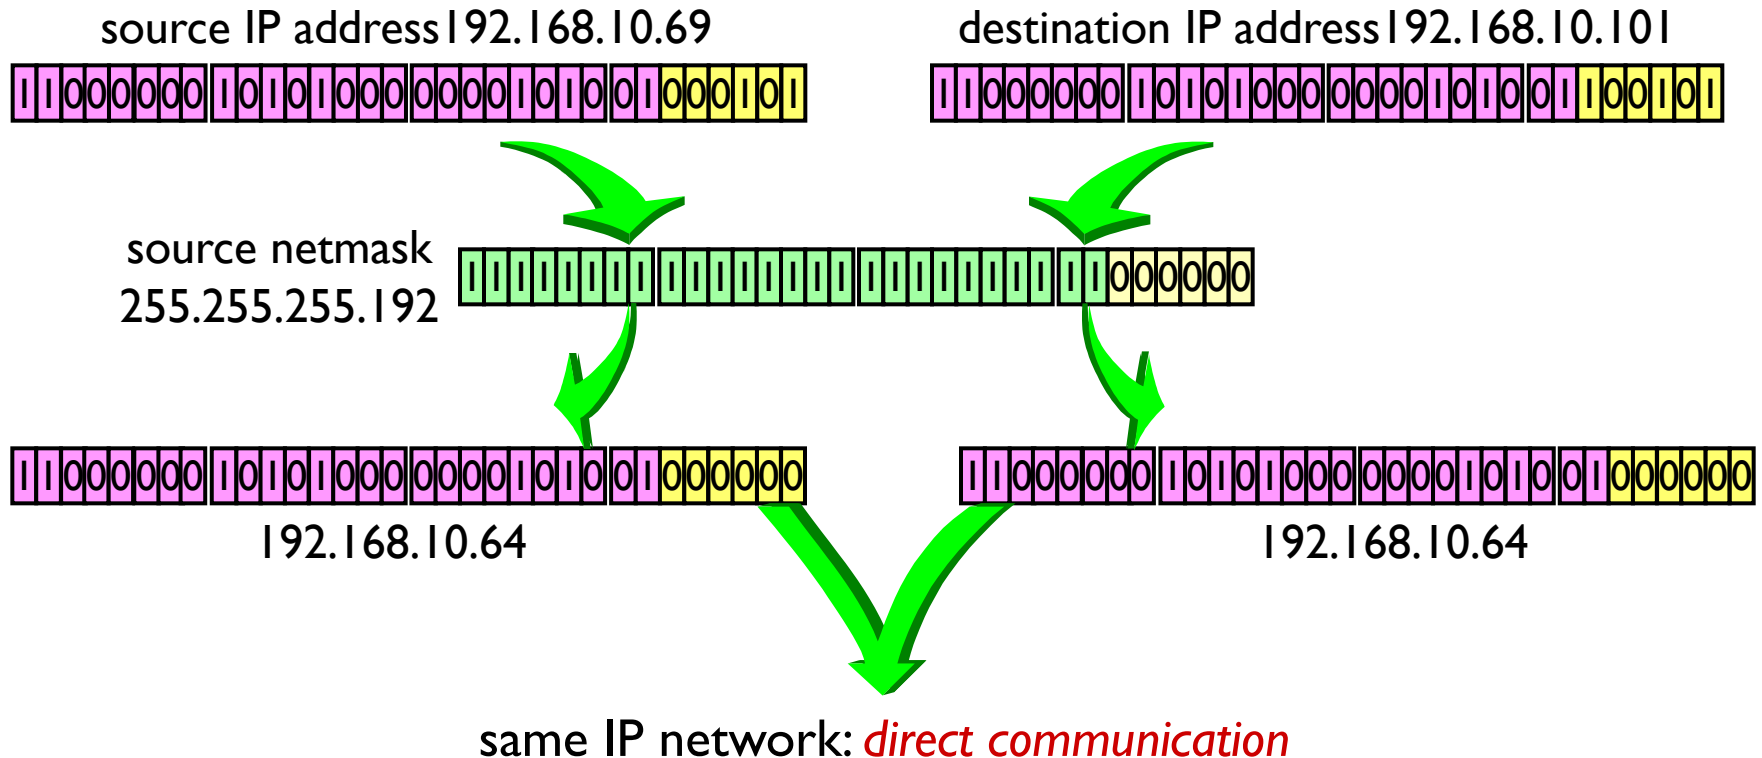
| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

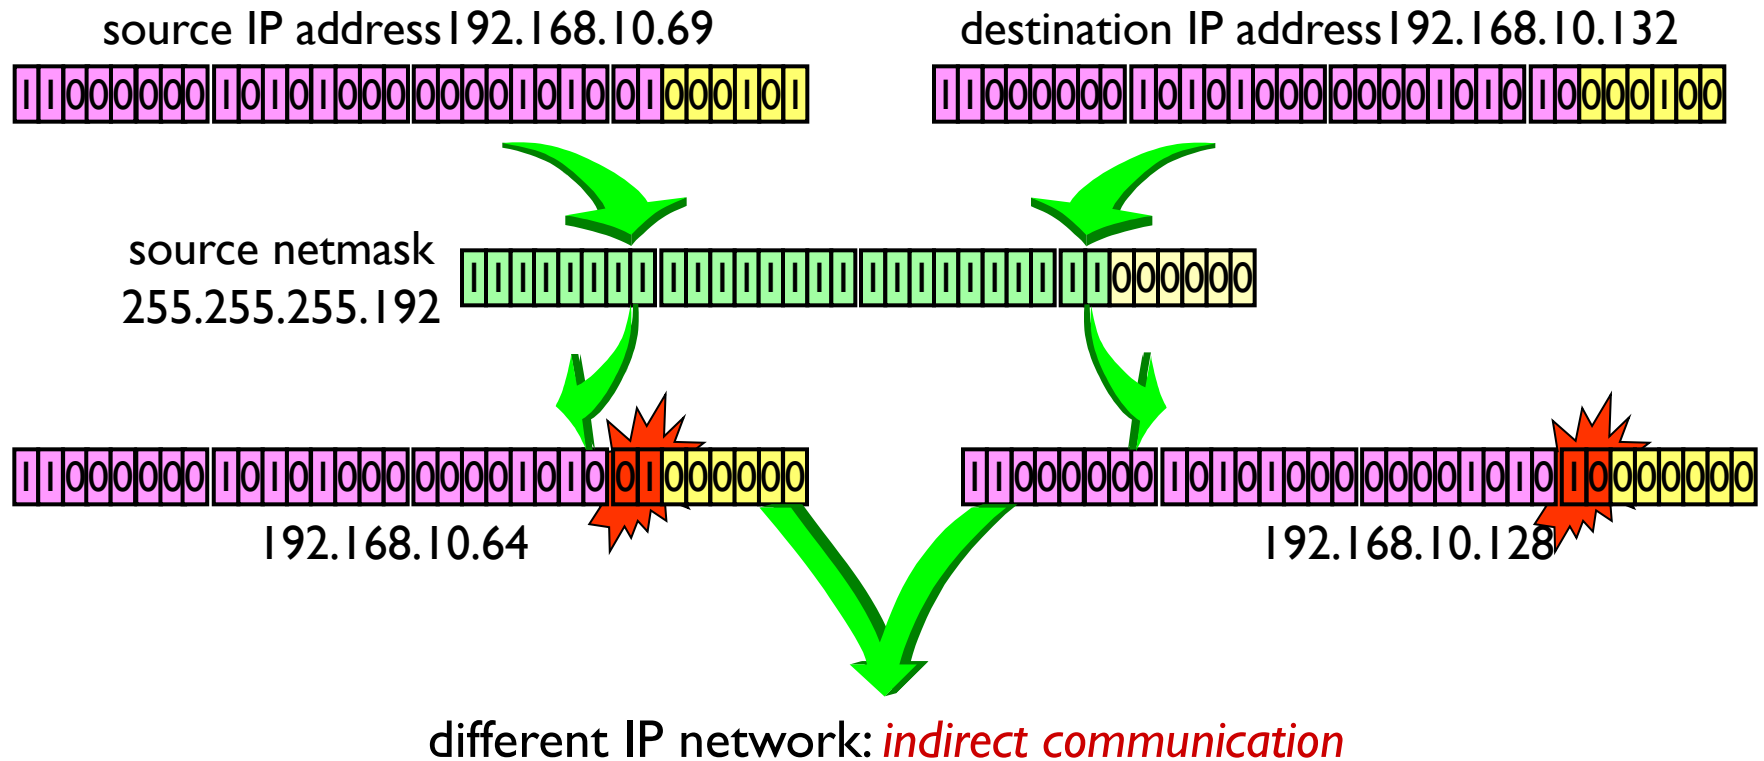
bit-wise AND result: 192.168.10.128

netmask: 255.255.255.192

How do IP routing really work?



How do IP routing really work?



we need a router!

How do IP routing really work?

Q: How do a router select the correct output port?

- ❖ *bit-wise AND* between the destination IP address of a packet and the netmask of each entry in the routing table, looking for a match

| routing table | |
|----------------|-------------|
| destination | output link |
| 200.23.16.0/20 | 1 |
| 199.31.0.0/16 | 2 |

How do IP routing really work?

Q: How do a router select the correct output port?

- ❖ *bit-wise AND* between the destination IP address of a packet and the netmask of each entry in the routing table, looking for a match

destination IP address: 200.23.16.1

11001000 00010111 00010000 00000001

11111111 11111111 11110000 00000000

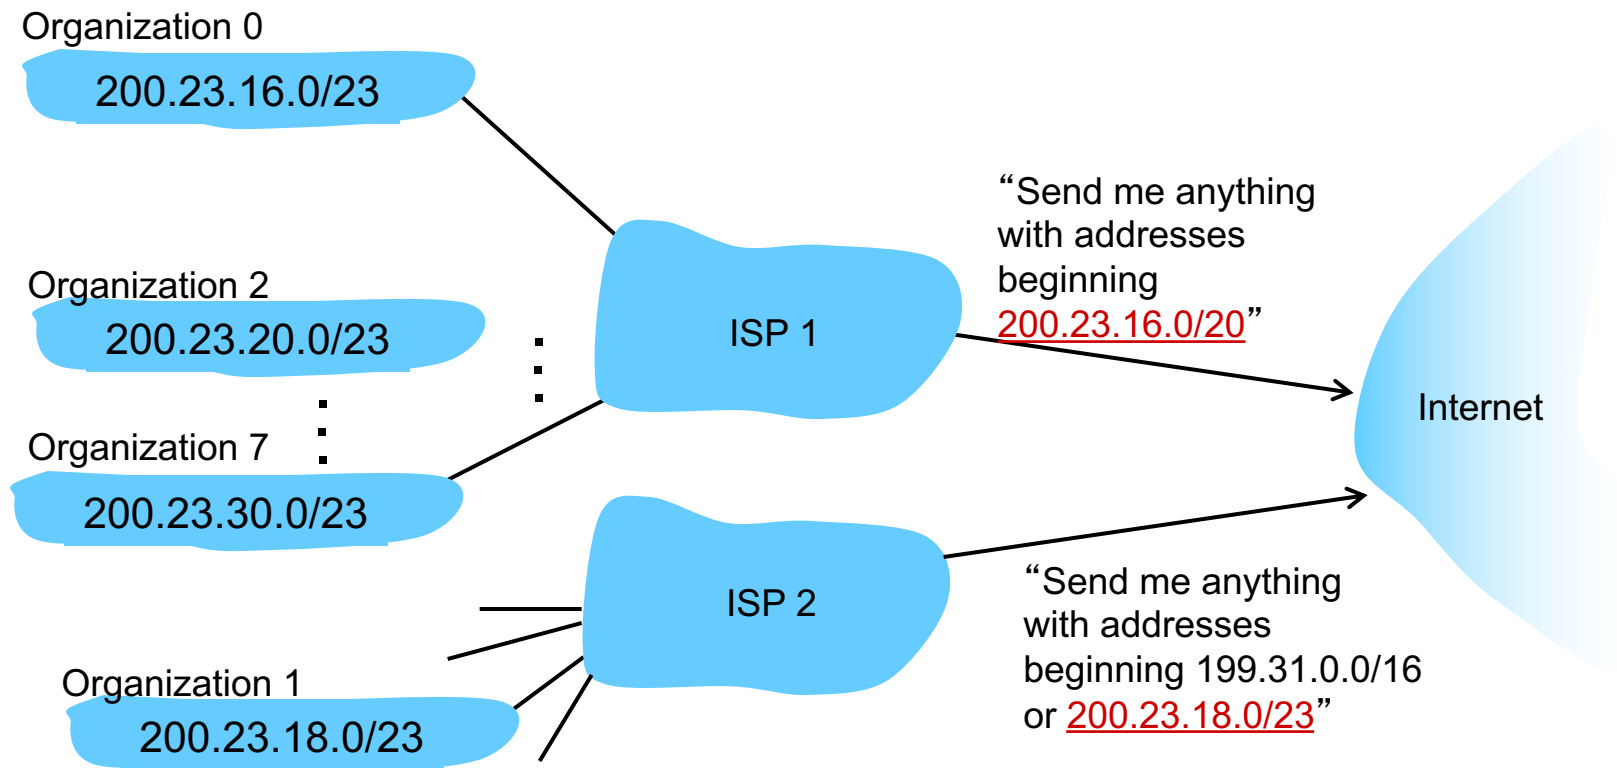
11001000 00010111 00010000 00000000

bit-wise AND result: 200.23.16.0 → match!

netmask: 255.255.240.0

Hierarchical addressing: more specific routes

ISP 2 has a more specific route to Organization 1



Longest prefix matching

longest prefix matching

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range | Link interface |
|-----------------------------|----------------|
| 11001000 00010111 0001***** | 0 |
| 11001000 00010111 0001001* | 1 |
| 11000111 00011111 ***** | 1 |
| otherwise | 2 |

examples:

DA: 11001000 00010111 00010110 10100001

which interface?

DA: 11001000 00010111 00010010 10101010

which interface?

IP routing: a real routing table

there are three types of *routes*

❖ *direct routes*

- networks directly connected to the router (i.e., address ranges including the router interfaces)

❖ *static routes*

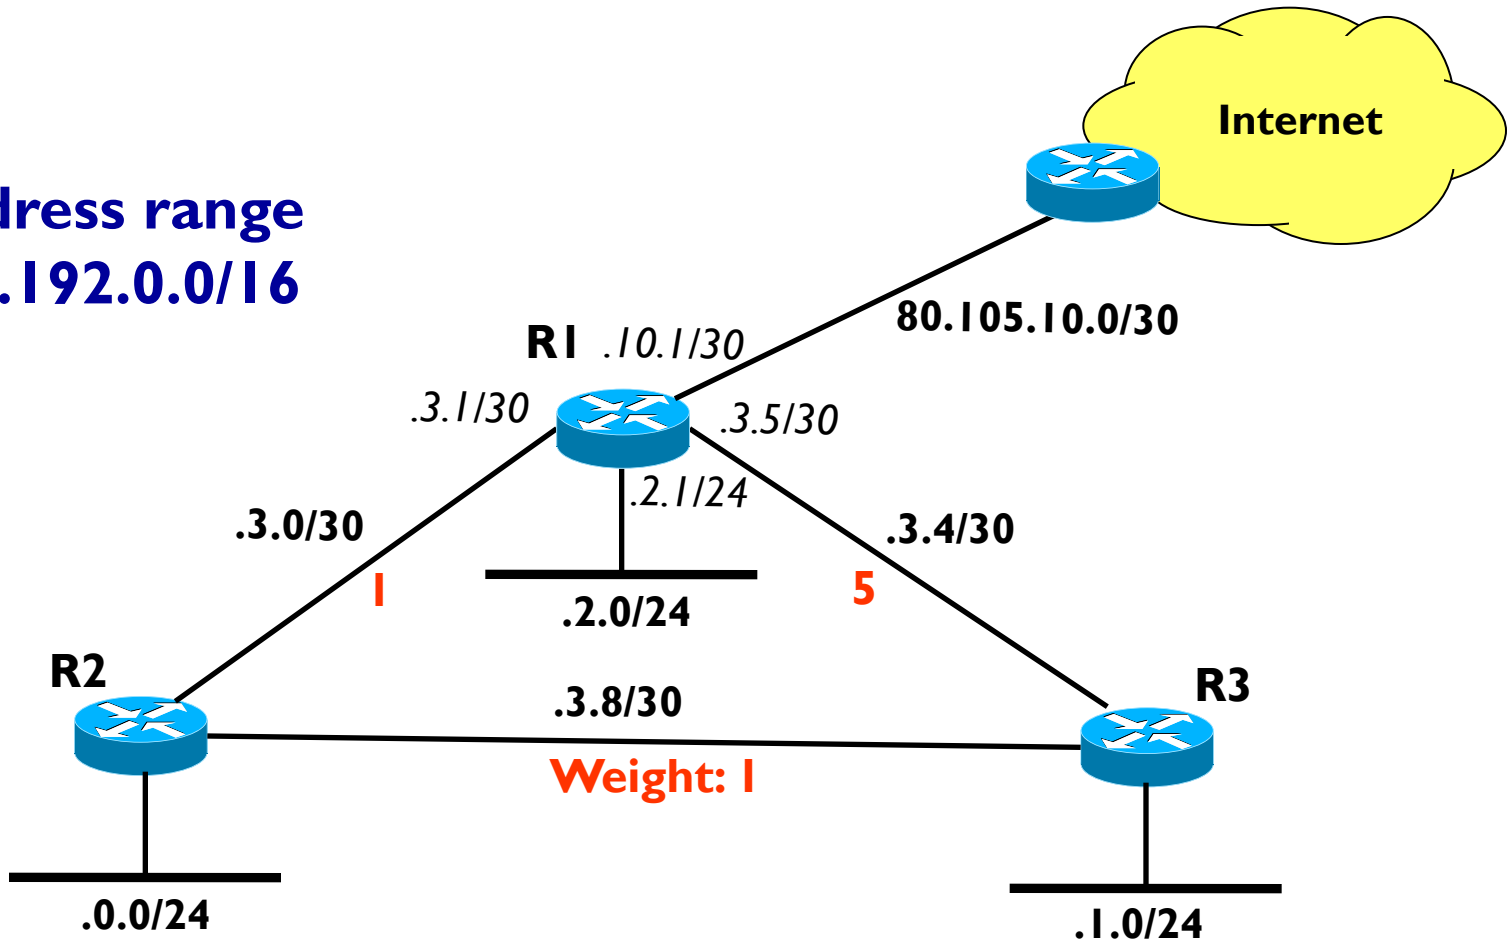
- manually configured routes to remote networks

❖ *dynamic routes*

- automatically configured routes to remote networks
 - routing protocols
 - ICMP redirect

IP routing: example

Address range
130.192.0.0/16



IP routing: example

RI routing table

| <i>Type</i> | <i>Destination</i> | <i>Next-hop</i> | <i>Cost</i> |
|-------------|--------------------|-----------------|-------------|
| S | 130.192.0.0/24 | 130.192.3.2 | 2 |
| S | 130.192.1.0/24 | 130.192.3.2 | 2 |
| S | 130.192.3.8/30 | 130.192.3.2 | 2 |
| S | 0.0.0.0/0 | 80.105.10.2 | 2 |
| D | 130.192.2.0/24 | 130.192.2.1 | 1 |
| D | 130.192.3.0/30 | 130.192.3.1 | 1 |
| D | 130.192.3.4/30 | 130.192.3.5 | 1 |
| D | 80.105.10.0/30 | 80.105.0.1 | 1 |

remote
interface!!

local
interface!!

IP addresses: how to get one?

Q: How does a *host* get IP address?

- ❖ hard-coded by system admin in a file
 - Windows: control-panel->network->configuration->tcp/ip->properties
 - UNIX: /etc/rc.config
- ❖ **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
 - “plug-and-play”

DHCP: Dynamic Host Configuration Protocol


goal: allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

DHCP overview:

- host broadcasts “DHCP discover” msg [optional]
- DHCP server responds with “DHCP offer” msg [optional]
- host requests IP address: “DHCP request” msg
- DHCP server sends address: “DHCP ack” msg

DHCP client-server scenario



DHCP client-server scenario

DHCP server: 223.1.2.5

DHCP discover

src : 0.0.0.0, 68
dest.: 255.255.255.255,67
yiaddr: 0.0.0.0
transaction ID: 654

arriving
client



DHCP offer

src: 223.1.2.5, 67
dest: 255.255.255.255, 68
yiaddr: 223.1.2.4
transaction ID: 654
lifetime: 3600 secs

DHCP request

src: 0.0.0.0, 68
dest.: 255.255.255.255, 67
yiaddr: 223.1.2.4
transaction ID: 655
lifetime: 3600 secs

DHCP ACK


src: 223.1.2.5, 67
dest: 255.255.255.255, 68
yiaddr: 223.1.2.4
transaction ID: 655
lifetime: 3600 secs

DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:


- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

DHCP: example



- ❖ connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- ❖ DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- ❖ Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server
- ❖ Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

DHCP: example



- ❖ DCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- ❖ encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- ❖ client now knows its IP address, name and IP address of DSN server, IP address of its first-hop router

DHCP: Wireshark output (home LAN)

Message type: **Boot Request (1)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

Transaction ID: 0x6b3a11b7

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

Client IP address: 0.0.0.0 (0.0.0.0)

Your (client) IP address: 0.0.0.0 (0.0.0.0)

Next server IP address: 0.0.0.0 (0.0.0.0)

Relay agent IP address: 0.0.0.0 (0.0.0.0)

Client MAC address: Wistron_23:68:8a (00:16:d3:23:68:8a)

Server host name not given

Boot file name not given

Magic cookie: (OK)

Option: (t=53,l=1) **DHCP Message Type = DHCP Request**

Option: (61) Client identifier

Length: 7; Value: 010016D323688A;

Hardware type: Ethernet

Client MAC address: Wistron_23:68:8a (00:16:d3:23:68:8a)

Option: (t=50,l=4) Requested IP Address = 192.168.1.101

Option: (t=12,l=5) Host Name = "nomad"

Option: (55) Parameter Request List

Length: 11; Value: 010F03062C2E2F1F21F92B

1 = Subnet Mask; 15 = Domain Name

3 = Router; 6 = Domain Name Server

44 = NetBIOS over TCP/IP Name Server

.....

request

Message type: **Boot Reply (2)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

Transaction ID: 0x6b3a11b7

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

Client IP address: 192.168.1.101 (192.168.1.101)

Your (client) IP address: 0.0.0.0 (0.0.0.0)

Next server IP address: 192.168.1.1 (192.168.1.1)

Relay agent IP address: 0.0.0.0 (0.0.0.0)

Client MAC address: Wistron_23:68:8a (00:16:d3:23:68:8a)

Server host name not given

Boot file name not given

Magic cookie: (OK)

Option: (t=53,l=1) DHCP Message Type = DHCP ACK

Option: (t=54,l=4) Server Identifier = 192.168.1.1

Option: (t=1,l=4) Subnet Mask = 255.255.255.0

Option: (t=3,l=4) Router = 192.168.1.1

Option: (6) Domain Name Server

Length: 12; Value: 445747E2445749F244574092;

IP Address: 68.87.71.226;

IP Address: 68.87.73.242;

IP Address: 68.87.64.146

Option: (t=15,l=20) Domain Name = "hsd1.ma.comcast.net."

reply

IP addressing: the last word...

Q: how does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes

NAT: network address translation

- ❖ some address ranges are reserved for private networks
 - they are not announced over the Internet, so they are not reachable from remote areas of the network
 - reserved as classful networks, but still valid in CIDR

10.0.0.0 - 10.255.255.255

1 class A network


172.16.0.0 - 172.31.255.255

16 class B networks

192.168.0.0 - 192.168.255.255

256 class C networks

NAT: network address translation



all datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0.0/24 address for source, destination (as usual)

NAT: network address translation

motivation: local network uses just one IP address as far as outside world is concerned:

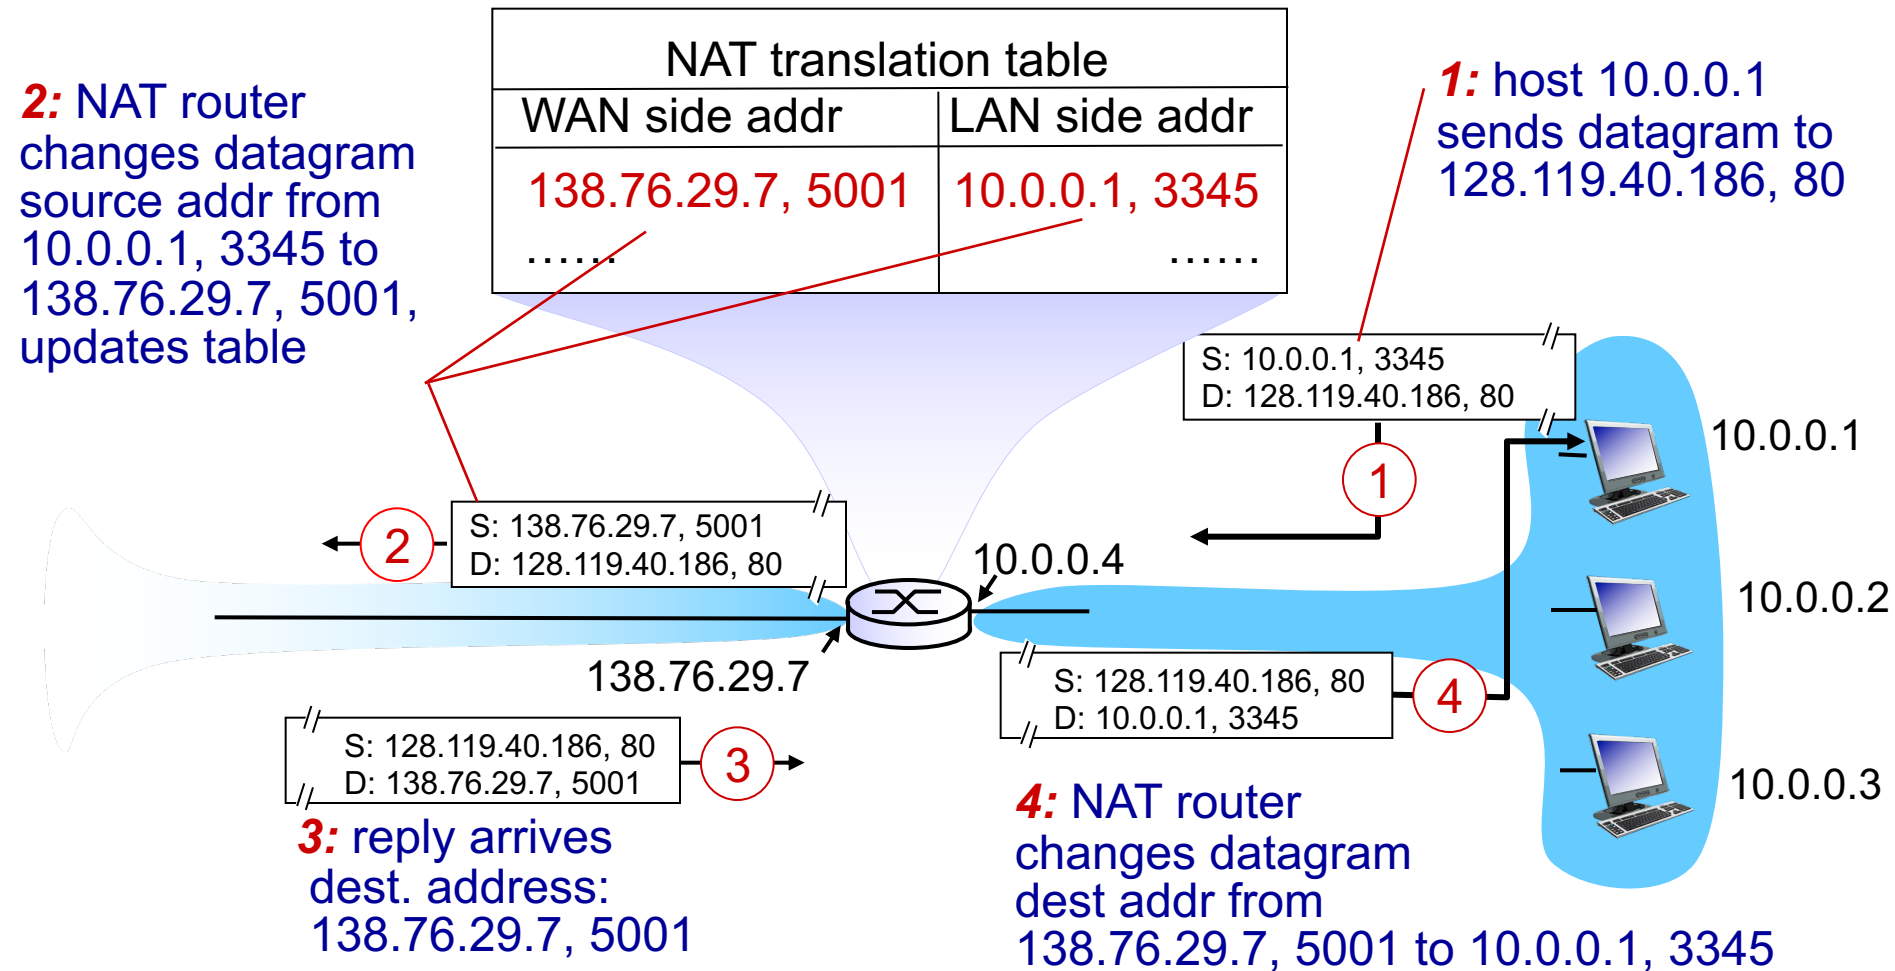
- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

NAT: network address translation

implementation: NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: network address translation




NAT: network address translation

- ❖ 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- ❖ NAT is controversial:
 - routers should only process up to layer 3
 - violates end-to-end argument
 - NAT possibility must be taken into account by app designers, e.g., P2P applications
 - address shortage should instead be solved by IPv6

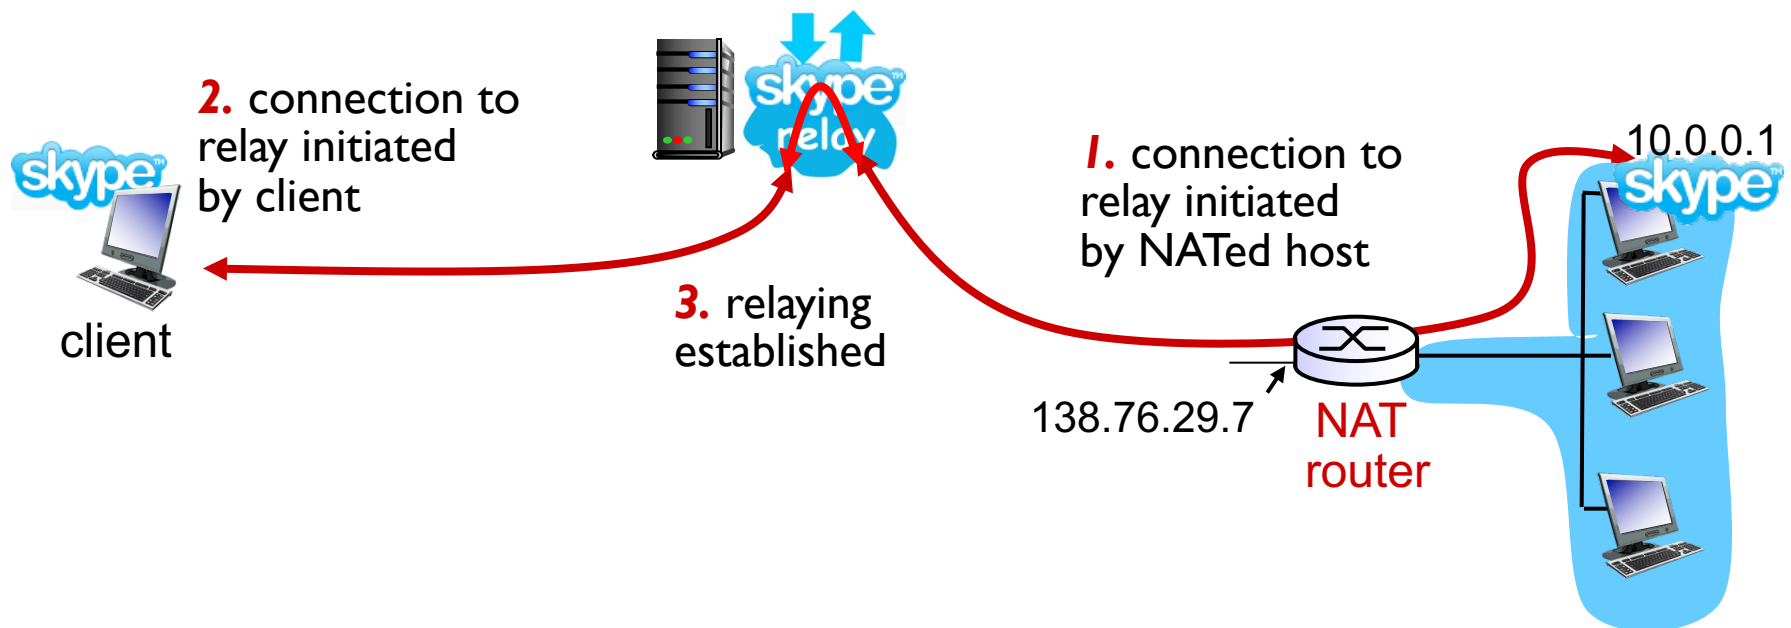
NAT traversal problem

- ❖ client wants to connect to server with address 10.0.0.1
 - server address 10.0.0.1 local to LAN (client can't use it as destination addr)
 - only one externally visible NATed address: 138.76.29.7
- ❖ **solution 1:** statically configure NAT to forward incoming connection requests at given port to server
 - e.g., (138.76.29.7, port 2500) always forwarded to 10.0.0.1 port 25000



NAT traversal problem

- ❖ **solution 2:** relaying (used in Skype)
 - NATed client establishes connection to relay
 - external client connects to relay
 - relay bridges packets between to connections



Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

- RIP
- OSPF
- BGP

ICMP: internet control message protocol

- ❖ used by hosts & routers to communicate network-level information

- error reporting:
unreachable host, network, port, protocol
- echo request/reply (used by ping)

- ❖ network-layer “above” IP:


- ICMP msgs carried in IP datagrams

- ❖ **ICMP message:** type, code plus first 8 bytes of IP datagram causing error

| <u>Type</u> | <u>Code</u> | <u>description</u> |
|-------------|-------------|---|
| 0 | 0 | echo reply (ping) |
| 3 | 0 | dest. network unreachable |
| 3 | 1 | dest host unreachable |
| 3 | 2 | dest protocol unreachable |
| 3 | 3 | dest port unreachable |
| 3 | 6 | dest network unknown |
| 3 | 7 | dest host unknown |
| 4 | 0 | source quench (congestion control - not used) |
| 5 | 0 | redirect |
| 8 | 0 | echo request (ping) |
| 9 | 0 | route advertisement |
| 10 | 0 | router discovery |
| 11 | 0 | TTL expired (time exceeded) |
| 12 | 0 | bad IP header |

Traceroute and ICMP

- ❖ source sends series of UDP segments to dest
 - first set has TTL = 1
 - second set has TTL=2, etc.
 - unlikely port number
 - ❖ when n th set of datagrams arrives to n th router:
 - router discards datagrams
 - and sends source ICMP messages (type 11, code 0)
 - ICMP messages includes name of router & IP address
 - ❖ when ICMP messages arrives, source records RTTs
- stopping criteria:*
- ❖ UDP segment eventually arrives at destination host
 - ❖ destination returns ICMP “port unreachable” message (type 3, code 3)
 - ❖ source stops



IPv6: motivation

- ❖ *initial motivation*: 32-bit address space soon to be completely allocated.
- ❖ additional motivation:
 - header format helps speed processing/forwarding
 - header changes to facilitate QoS

IPv6 datagram format:

- fixed-length 40 byte header
- no fragmentation allowed


IPv6 datagram format

priority: identify priority among datagrams in flow

flow Label: identify datagrams in same “flow.”

(concept of “flow” not well defined).

next header: identify upper layer protocol for data

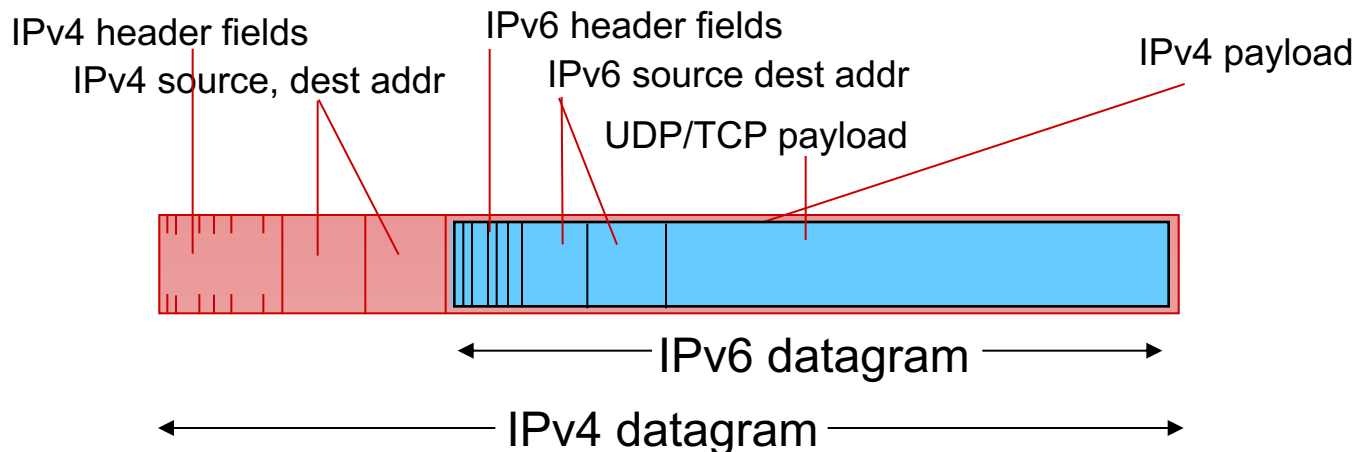


Other changes from IPv4

- ❖ *checksum*: removed entirely to reduce processing time at each hop
- ❖ *options*: allowed, but outside of header, indicated by “Next Header” field
- ❖ *ICMPv6*: new version of ICMP
 - additional message types, e.g. “Packet Too Big”
 - multicast group management functions


Transition from IPv4 to IPv6

- ❖ not all routers can be upgraded simultaneously
 - no “flag days”
 - how will network operate with mixed IPv4 and IPv6 routers?
- ❖ **tunneling**: IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers




Tunneling

logical view:




physical view:





Tunneling

logical view:



physical view:



Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

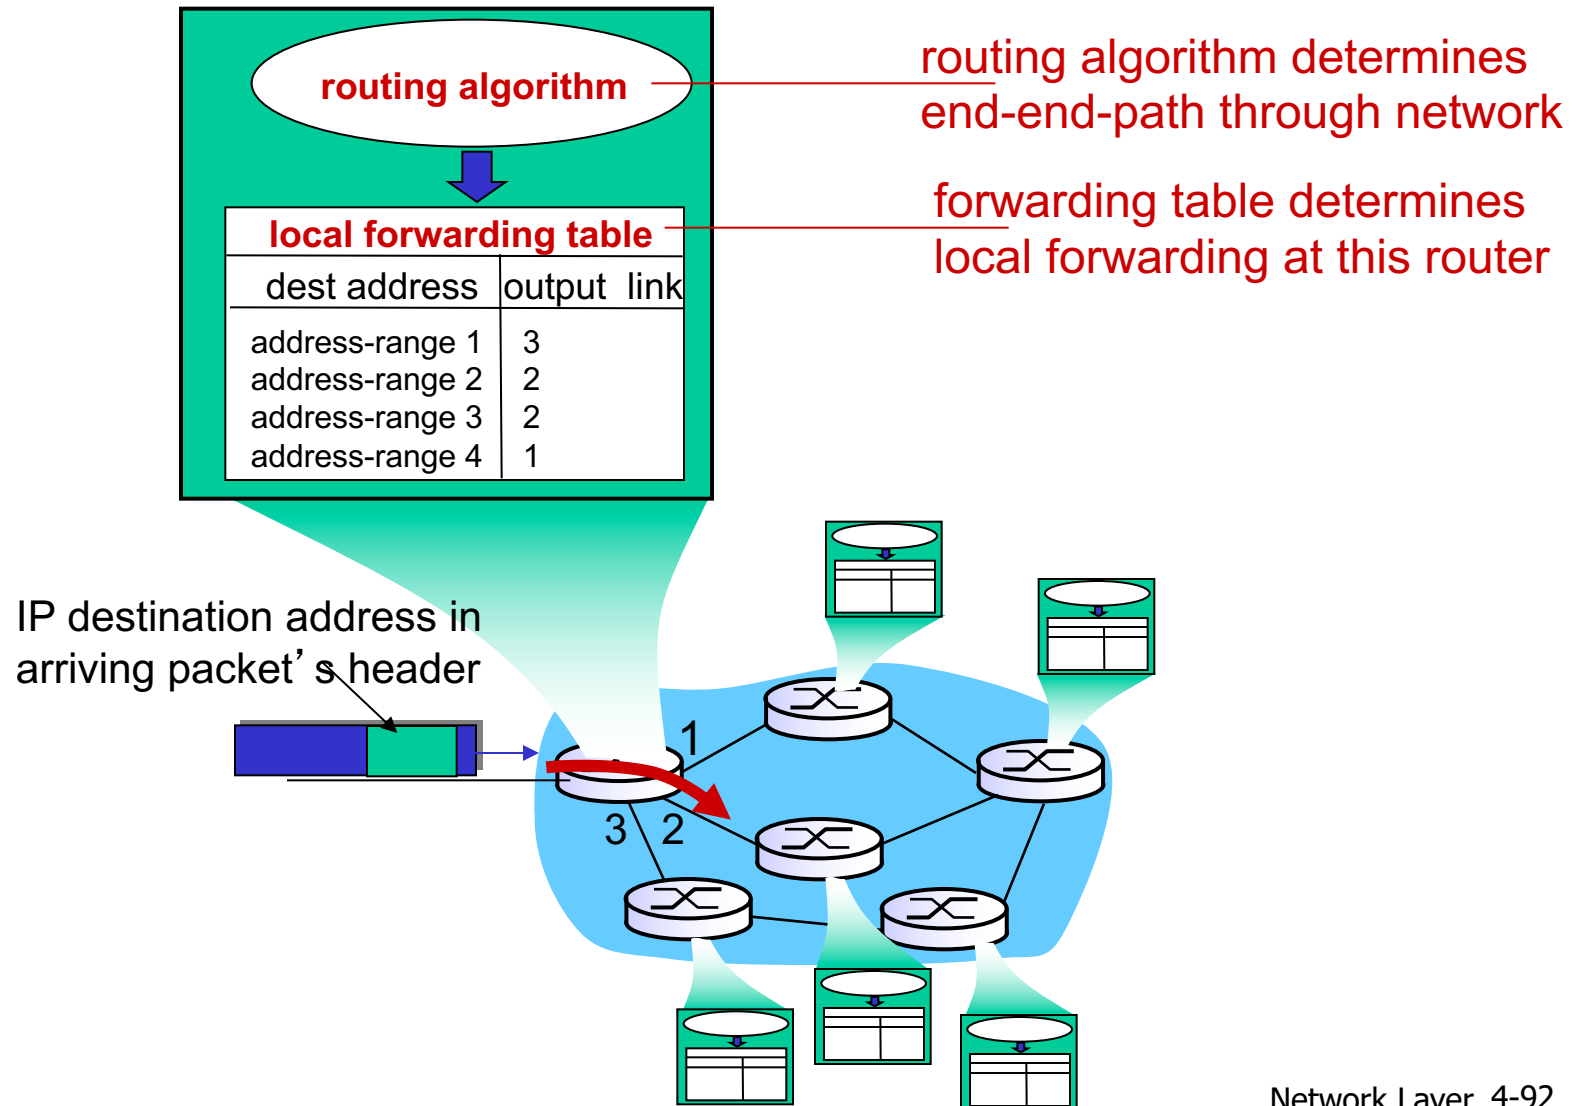
4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

- RIP
- OSPF
- BGP

Interplay between routing, forwarding



Comparison of LS and DV algorithms

message complexity

- ❖ **LS:** with n nodes, E links, $O(nE)$ msgs sent
- ❖ **DV:** exchange between neighbors only
 - convergence time varies

speed of convergence

- ❖ **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- ❖ **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its own table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet

- RIP
- OSPF
- BGP

Hierarchical routing

our routing study thus far - idealization

- ❖ all routers identical

- ❖ network “flat”

... *not* true in practice

scale: with 600 million destinations:

- ❖ can't store all dest's in routing tables!
- ❖ routing table exchange would swamp links!

administrative autonomy

- ❖ internet = network of networks
- ❖ each network admin may want to control routing in its own network

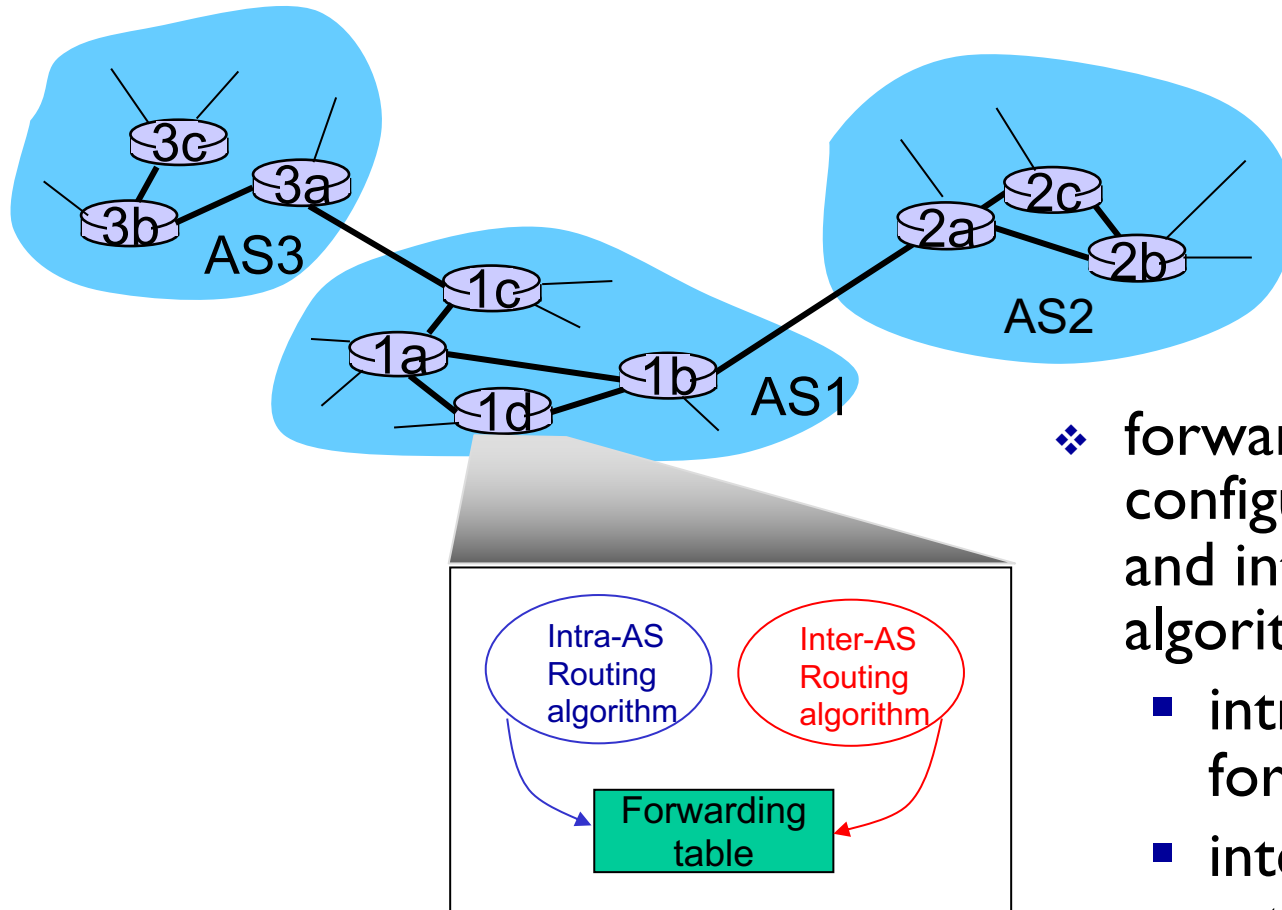
Hierarchical routing

- ❖ aggregate routers into regions, “**autonomous systems**” (AS)
- ❖ routers in same AS run same routing protocol
 - “**intra-AS**” routing protocol
 - routers in different AS can run different intra-AS routing protocol

gateway router:

- ❖ at “edge” of its own AS
- ❖ has link to router in another AS

Interconnected ASes



- ❖ forwarding table configured by both intra- and inter-AS routing algorithm
 - intra-AS sets entries for internal dests
 - inter-AS & intra-AS sets entries for external dests


Inter-AS tasks

- ❖ suppose router in AS1 receives datagram destined outside of AS1:
 - router should forward packet to gateway router, but which one?

AS1 must:


1. learn which destds are reachable through AS2, which through AS3
2. propagate this reachability info to all routers in AS1

job of inter-AS routing!




Example: setting forwarding table in router 1d

- ❖ suppose AS1 learns (via inter-AS protocol) that subnet **x** reachable via AS3 (gateway 1c), but not via AS2
 - inter-AS protocol propagates reachability info to all internal routers
- ❖ router 1d determines from intra-AS routing info that its interface **l** is on the least cost path to 1c
 - installs forwarding table entry **(x,l)**



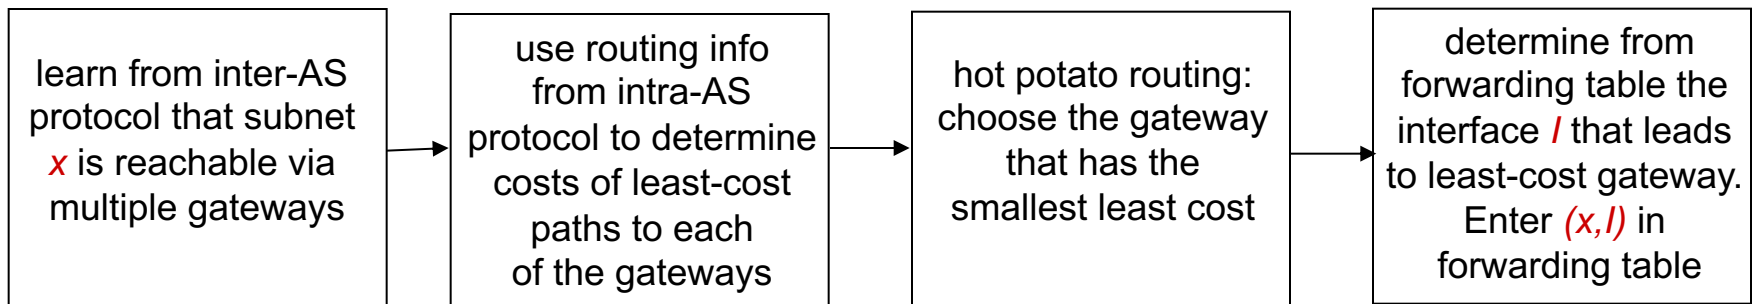
Example: choosing among multiple ASes

- ❖ now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ❖ to configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest **x**
 - this is also job of inter-AS routing protocol!



Example: choosing among multiple ASes

- ❖ now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ❖ to configure forwarding table, router 1d must determine towards which gateway it should forward packets for dest **x**
 - this is also job of inter-AS routing protocol!
- ❖ **hot potato routing: send** packet towards closest of two routers.



Chapter 4: outline

4.1 introduction

4.2 virtual circuit and datagram networks

4.3 IP: Internet Protocol

- datagram format
- IPv4 addressing
- ICMP
- IPv6

4.4 routing algorithms

- link state
- distance vector
- hierarchical routing

4.5 routing in the Internet


- RIP
- OSPF
- BGP

Intra-AS Routing

- ❖ also known as *interior gateway protocols (IGP)*
- ❖ most common intra-AS routing protocols:
 - RIP: Routing Information Protocol
 - OSPF: Open Shortest Path First
 - IGRP: Interior Gateway Routing Protocol (Cisco proprietary)

RIP (Routing Information Protocol)

- ❖ included in BSD-UNIX distribution in 1982
- ❖ distance vector algorithm
 - distance metric: # hops (max = 15 hops), each link has cost 1
 - DVs exchanged with neighbors every 30 sec in response message (aka **advertisement**)
 - each advertisement: list of up to 25 destination **subnets** (in IP addressing sense)




from router A to destination **subnets**:

| <u>subnet</u> | <u>hops</u> |
|---------------|-------------|
| u | 1 |
| v | 2 |
| w | 2 |
| x | 3 |
| y | 3 |
| z | 2 |

OSPF (Open Shortest Path First)

- ❖ “open”: publicly available
- ❖ uses link state algorithm
 - LS packet dissemination
 - topology map at each node
 - route computation using Dijkstra’s algorithm
- ❖ OSPF advertisement carries one entry per neighbor
- ❖ advertisements flooded to *entire* AS
 - carried in OSPF messages directly over IP (rather than TCP or UDP)
- ❖ *IS-IS routing* protocol: nearly identical to OSPF

Hierarchical OSPF



Internet inter-AS routing: BGP

- ❖ **BGP (Border Gateway Protocol):** *the de facto inter-domain routing protocol*
 - “glue that holds the Internet together”
- ❖ BGP provides each AS a means to:
 - **eBGP:** obtain subnet reachability information from neighboring ASs.
 - **iBGP:** propagate reachability information to all AS-internal routers.
 - determine “good” routes to other networks based on reachability information and policy.
- ❖ allows subnet to advertise its existence to rest of Internet: *“I am here”*

Why different Intra-, Inter-AS routing ?

policy:

- ❖ inter-AS: admin wants control over how its traffic routed, who routes through its net.
- ❖ intra-AS: single admin, so no policy decisions needed

scale:

- ❖ hierarchical routing saves table size, reduced update traffic

performance:

- ❖ intra-AS: can focus on performance
- ❖ inter-AS: policy may dominate over performance